

PRIYANSHU RAJ

9006720546 | ipriyanshu009@gmail.com | github.com/priyanshuraj009 | linkedin.com/in/priyanshuraj009

PROFESSIONAL SUMMARY

Generative AI and RAG specialist with 1.5+ years at NVIDIA (via Randstad), building and evaluating LLM-powered pipelines for large-scale autonomous vehicle and computer vision AI programs. Hands-on with RAG pipeline design, Lang Chain and Lang Graph orchestration, and multi-model prompt engineering across GPT-4, Claude, and Gemini. Delivered 90%+ quality outcomes across Dense Captioning, COSMOS AV, Temporal Reasoning, Skeleton Motion Description, and Human Egocentric projects. Strong analytical foundation in data pre-processing, statistical validation, and structured dataset QA, enabling data-informed decisions that improve model performance and reliability.

WORK EXPERIENCE

Associate Prompt Engineer

October 2024 - May 2026

NVIDIA Graphics - Contract via Randstad | Remote

- Designed RAG-style prompt pipelines for Dense Captioning and Temporal Reasoning projects, integrating contextual retrieval and Chain-of-Thought reasoning to improve LLM factual grounding, reduce hallucinations, and enhance structured output quality across production AI systems.
- Developed and iterated few-shot and instruction-tuned prompt strategies using GPT-4, Claude, and Gemini across five NVIDIA AI projects; maintained 90%+ quality while managing full QC lifecycle for Skeleton Motion Description dataset (95%+ QA, 15,000+ annotations).
- Orchestrated Lang Chain and Lang Graph-based prompt workflows for Human Egocentric and AV Alpamayo perception tasks; applied statistical benchmarking and NLP-based output analysis to validate model-generated captions against ground-truth labels.
- Executed quality control for COSMOS AV generative dataset, refining prompt templates in collaboration with ML engineers; produced RLHF-aligned structured feedback that supported model fine-tuning and instruction-tuning pipelines.

TECHNICAL SKILLS

Generative AI and RAG: RAG Pipelines, Prompt Engineering, CoT Reasoning, Instruction Tuning, Few-Shot and Zero-Shot Prompting, RLHF, Fine-tuning

LLM and Frameworks: GPT-4, Claude (Anthropic), Gemini, Lang Chain, Lang Graph, OpenAI API

Data and Analysis: Statistical Analysis, Data Pre-processing, Dataset QA, Feature Engineering, NLP, Pandas, SQL

ML and Computer Vision: Scikit-learn, TensorFlow, PyTorch, Object Detection, Video Annotation, AV Perception

Programming and Tools: Python, SQL, Java, Git, Jupyter Notebook, VS Code, Twilio, Google Speech-to-Text

PROJECTS

Med-Voice - Python, Twilio, Google Speech-to-Text, NLP | Hackathon 2023

- Built a generative speech-to-text pipeline for clinical documentation using Twilio and Google Speech-to-Text APIs, transforming unstructured audio into structured medical records in real time and eliminating paper-based clinical workflows.

Fraud Detection ML Model - Python, Scikit-learn, Pandas | 2024

- Developed an ensemble fraud detection model (Random Forest) with SMOTE oversampling and feature engineering; applied precision-recall threshold tuning on imbalanced real-world financial datasets to optimize detection accuracy.

Intelligent Battery Management System - Research Publication | IJRAR 2024

- Published peer-reviewed ML research applying data-driven modelling to intelligent BMS optimization, demonstrating structured experimentation and applied AI research capabilities.

EDUCATION

B.Tech in Computer Science

September 2020 - June 2024

Presidency University, Bangalore | CGPA: 8.36 / 10.00

- Coursework: Data Science, Deep Learning, Reinforcement Learning, Database Management, Software Project Management. Design Head across multiple university clubs - led branding, event collateral, and digital creative direction.

CERTIFICATIONS

- Essentials of Prompt Engineering - AWS
- Machine Learning with Python Level 1 - IBM
- Data Mining - Great Learning
- Java, Python, MySQL - HackerRank